

Semantically Enhancing OLAP Cubes: Integrating SPARQL and SQL for Next-Generation Data Publication & Business Intelligence

Benjamin Cogrel and Adrian Gschwend

Knowledge Graph Forum, Basel, Switzerland, May 30th 2024



CSV, XLS, ...; let's call it table

	A	B	C	D
1	id	cityId	temp	timestamp
2	1	Q425415	-17	2009-05-19T06:32:08
3	2	Q671288	19.9	2003-11-26T11:33:18
4	3	Q956602	-26.9	2011-08-29T20:02:38
5	4	Q131638	44.8	2001-11-05T22:01:03
6	5	Q13025347	12	2011-12-03T08:09:07
7	6	Q10976	42.9	2006-03-10T05:17:37

	A	B	C
1	city	cityLabel	cityId
2	http://www.wikidata.org/entity/Q44214	Formosa	Q44214
3	http://www.wikidata.org/entity/Q41252	Bydgoszcz	Q41252
4	http://www.wikidata.org/entity/Q43433	Chandigarh	Q43433
5	http://www.wikidata.org/entity/Q38811	Nasiriyah	Q38811
6	http://www.wikidata.org/entity/Q42763	Santiago de los Caballeros	Q42763
7	http://www.wikidata.org/entity/Q43509	Guayaquil	Q43509
8	http://www.wikidata.org/entity/Q39984	Cannes	Q39984
9	http://www.wikidata.org/entity/Q40921	Honiara	Q40921
10	http://www.wikidata.org/entity/Q44162	San Fernando del Valle de Catamarca	Q44162
11	http://www.wikidata.org/entity/Q44237	Mendoza	Q44237
12	http://www.wikidata.org/entity/Q44239	Neuquén	Q44239

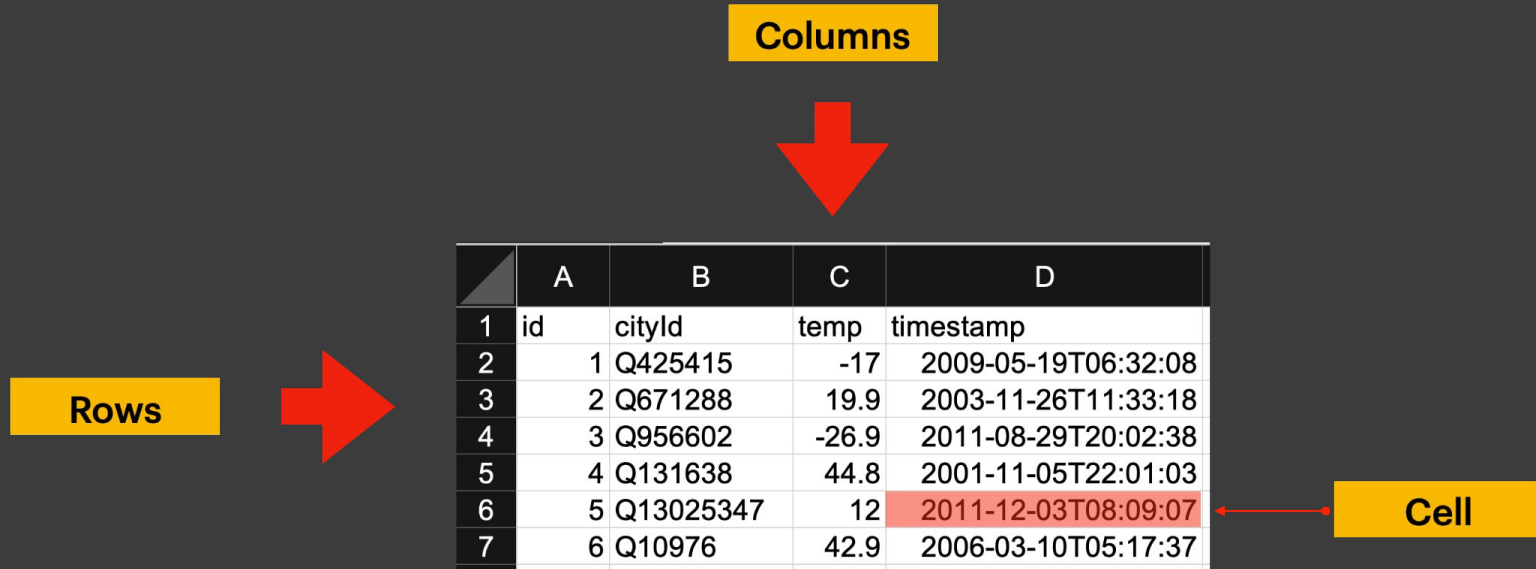


Tables

	A	B	C	D
1	id	cityId	temp	timestamp
2	1	Q425415	-17	2009-05-19T06:32:08
3	2	Q671288	19.9	2003-11-26T11:33:18
4	3	Q956602	-26.9	2011-08-29T20:02:38
5	4	Q131638	44.8	2001-11-05T22:01:03
6	5	Q13025347	12	2011-12-03T08:09:07
7	6	Q10976	42.9	2006-03-10T05:17:37

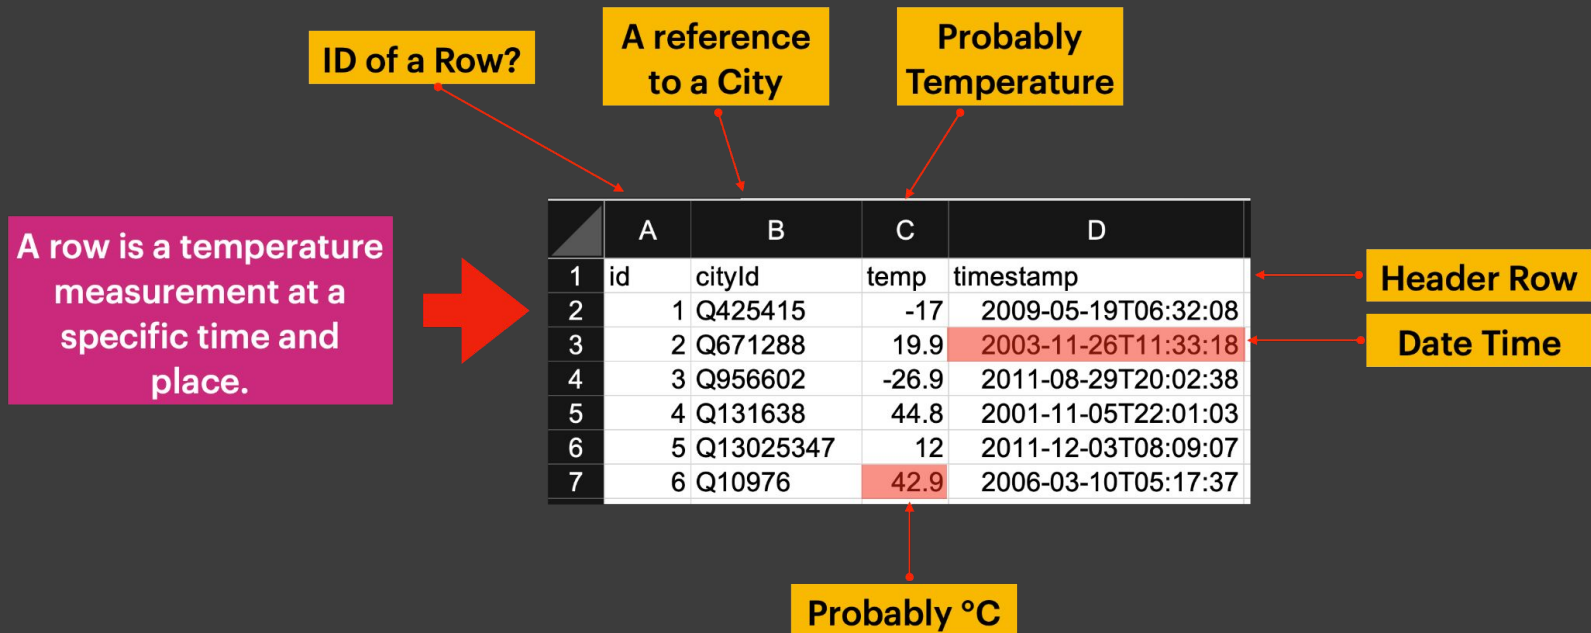


Semantics of Tables?



Implicit Semantics

Implicit semantics refer to the unspoken or inferred meaning that **you** derive from context.



The City Table

Implicit semantics refer to the unspoken or inferred meaning that **you** derive from context.

An URL called City ?

The name of the City

The ID of a City
Probably the same ID as in the previous Table

A row represents a city with its name and IDs

	A	B	C
1	city	cityLabel	cityId
2	http://www.wikidata.org/entity/Q44214	Formosa	Q44214
3	http://www.wikidata.org/entity/Q41252	Bydgoszcz	Q41252
4	http://www.wikidata.org/entity/Q43433	Chandigarh	Q43433
5	http://www.wikidata.org/entity/Q38811	Nasiriyah	Q38811
6	http://www.wikidata.org/entity/Q42763	Santiago de los Caballeros	Q42763
7	http://www.wikidata.org/entity/Q43509	Guayaquil	Q43509
8	http://www.wikidata.org/entity/Q39984	Cannes	Q39984
9	http://www.wikidata.org/entity/Q40921	Honiara	Q40921
10	http://www.wikidata.org/entity/Q44162	San Fernando del Valle de Catamarca	Q44162
11	http://www.wikidata.org/entity/Q44237	Mendoza	Q44237
12	http://www.wikidata.org/entity/Q44239	Neuquén	Q44239



Stop guessing semantics; make them explicit

Dimensions



Observation




	A	B	C	D
1	id	cityId	temp	timestamp
2	1	Q425415	-17	2009-05-19T06:32:08
3	2	Q671288	19.9	2003-11-26T11:33:18
4	3	Q956602	-26.9	2011-08-29T20:02:38
5	4	Q131638	44.8	2001-11-05T22:01:03
6	5	Q13025347	12	2011-12-03T08:09:07
7	6	Q10976	42.9	2006-03-10T05:17:37

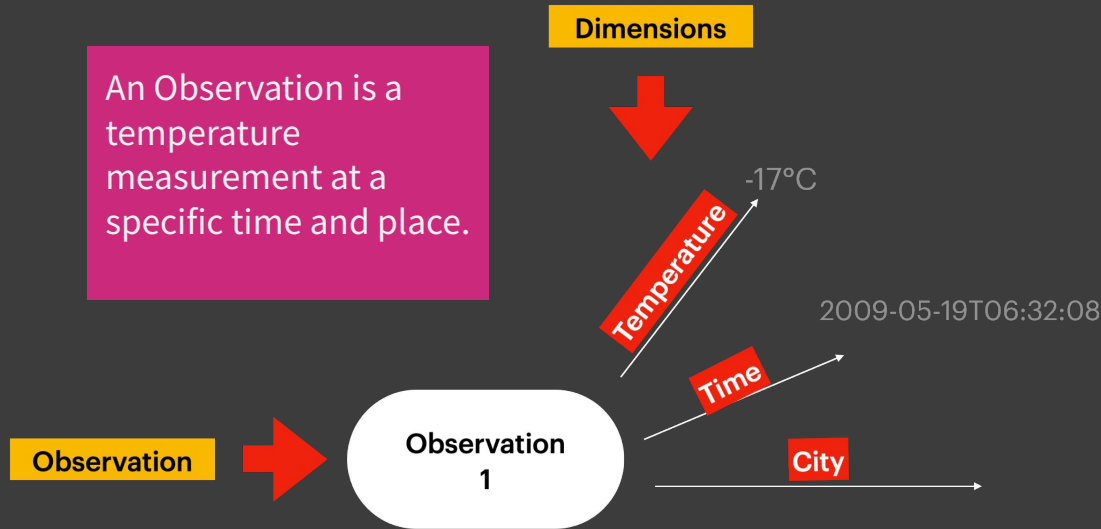


Stop guessing semantics; make them explicit

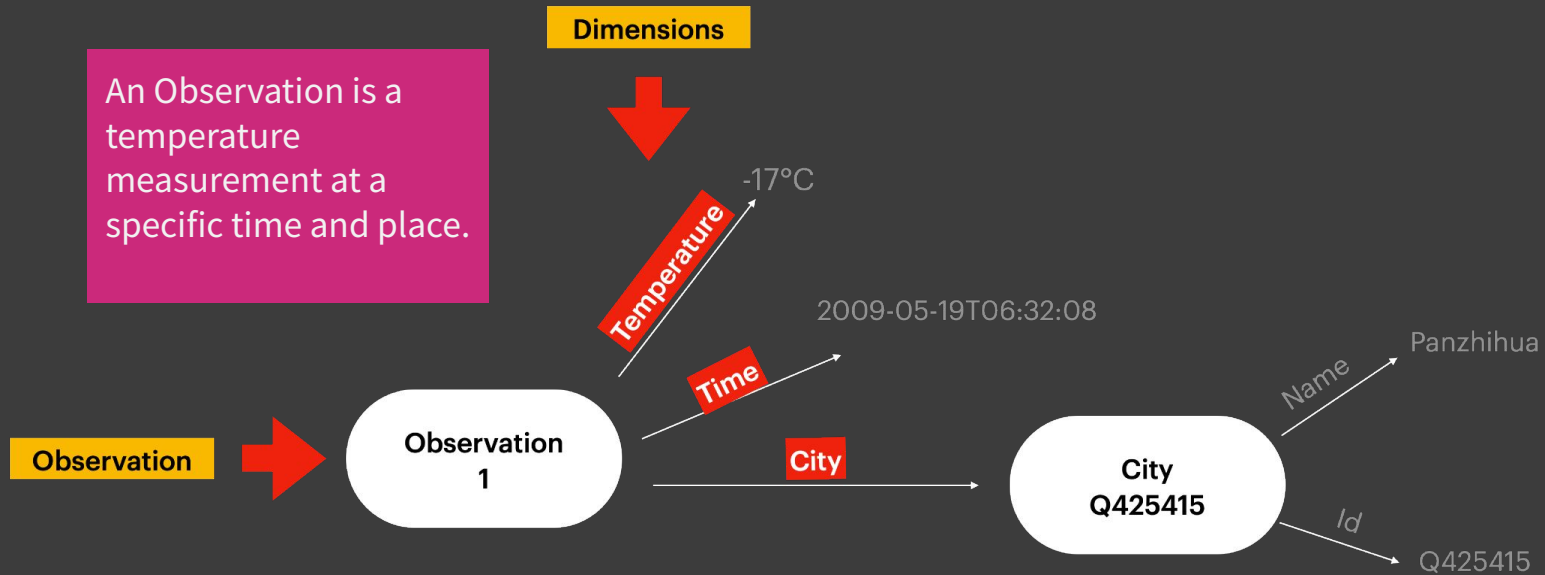
An Observation is a temperature measurement at a specific time and place.



Stop guessing semantics; make them explicit



Stop guessing semantics; make them explicit



Stop guessing semantics; make them explicit

The city where it was measured

Temperature in °C
between -65 and 55.

Time when it was
measured between
2009 and 2020

Dimension Metadata

- Name
- Description
- Unit
- Levels (Scale)
- Datatype
- ...

Cube Metadata

- Name
- Date
- Source
- ...

Observations



	A	B	C	D
1	id	cityId	temp	timestamp
2	1	Q425415	-17	2009-05-19T06:32:08
3	2	Q671288	19.9	2003-11-26T11:33:18
4	3	Q956602	-26.9	2011-08-29T20:02:38
5	4	Q131638	44.8	2001-11-05T22:01:03
6	5	Q13025347	12	2011-12-03T08:09:07
7	6	Q10976	42.9	2006-03-10T05:17:37



All this is a (OLAP) Cube

Online Analytical Processing with a multi-dimensional array (hypercube)



RDF Cube Vocabularies

- Cube.link
 - Swiss government/Zazuko
 - used in all the demos shown before
- schema.org
 - schema:Observation
- Semantic Sensor Network Ontology
 - sosa:Observation
- RDF Data Cube Vocabulary
 - The original

“Observation” is the common theme, details differ

But Why?


ZAZUKO CUBE VIEWER

SPARQL endpoint URL

Username


Password

Source graph





But Why?

Schweizerische Eidgenossenschaft
Confédération suisse
Confederazione Svizzera
Confederaziun svizra

visualize.admin.ch
Data source: [Int](#) [↗](#)




DE FR IT **EN**
[Sign in](#)

Visualize Swiss Open Government Data

Create and embed visualizations from any dataset provided by the LINDAS Linked Data Service.

Create a visualization

Visualize data in just a few steps...



Select a datasetEdit the visualizationShare & embed

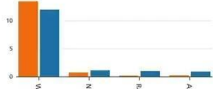
Make it your own...

Traffic noise pollution

Metadata

Show Filters +

Percentage (%)

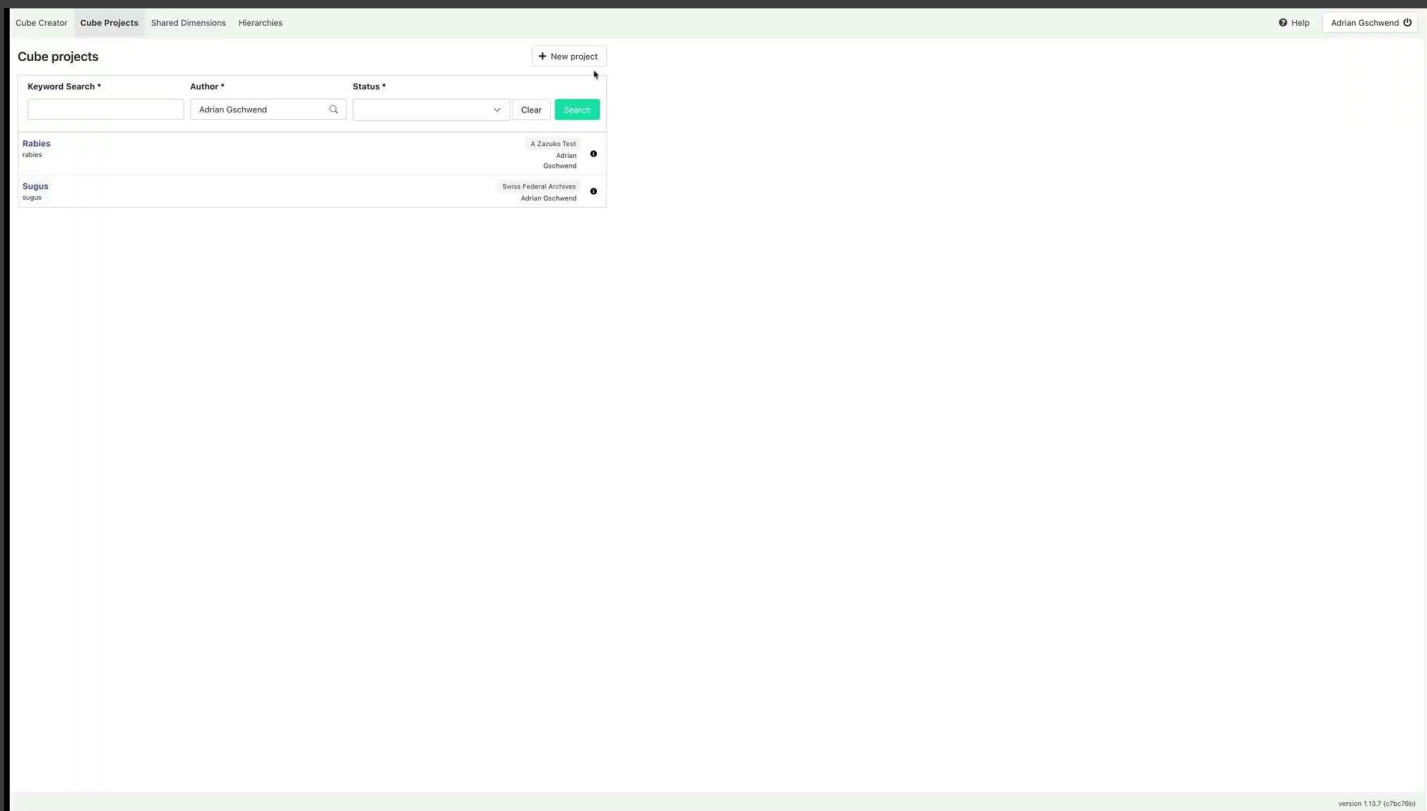


Category	Percentage (%)
Noise	10
Noise	10
Noise	1
Noise	1
Noise	1
Noise	1
Noise	1
Noise	1
Noise	1
Noise	1

Create beautiful visualizations



How?



RDF cubes with virtualization

- If your data in CSV/JSON/Parquet files or in a relational database
 - You can map it to RDF using an R2RML mapping
 - Systems like Dremio, Trino and DuckDB allow you to query files in SQL
- R2RML mappings enable 2 kinds of deployments
 - Materialization: data is transformed in RDF and loaded in a triplestore
 - Virtualization: user queries are translated into SQL and send to the data source
- 2 languages to query the KG
 - SPARQL
 - Ontop if virtual, triplestore otherwise
 - SQL (new!)
 - For Virtual KGs only (as of now)

Querying KGs in SQL

- All the default graph data is automatically available through virtual relational tables
 - Don't have to create custom SQL views (using a SPARQL query)
 - No extra-modelling effort, on a par with the SPARQL interface
 - One table per class, one per object property, one per data property
- Virtual tables contain as many columns as possible
 - They are denormalized so as to save users from writing joins
 - While still guaranteeing that there is no more than 1 row per entity
 - Safe to run aggregates over
 - Made possible by analyzing the structures of the mapping and of the source
 - Backed by many optimizations to run fast
- With foreign keys to pursue the linked data experience
 - For the 1-1 and n-1 relationships

Database Navigator x Projects

Enter a part of object name here

- Cube-ODH - ostest2.ontopic.dev:4300
 - Databases
 - Cube-ODH
 - Schemas
 - classes
 - Tables
 - Constraint**
 - Cube
 - Datatype
 - KeyDimension
 - MeasureDimension
 - Observation
 - ObservationSet
 - PropertyShape
 - Resource
 - Shape

Constraint x

Properties Data ER Diagram

Constraint Enter a SQL expression to filter results (use Ctrl+Space)

Grid	Row #228
constraint_id	http://odh.example.org/constraint/23/8
measureDimension	http://odh.example.org/measurement-dimension/8
measureDimension.description	Temperatura dell'aria
measureDimension.name	air-temperature
measureDimension.datatype	http://www.w3.org/2001/XMLSchema#decimal
measureDimension.path	http://vocabulary.example.org/airTemperature
sensorDimension	http://odh.example.org/sensor-dimension/23/8
sensorDimension.in	http://odh.example.org/sensor-list/23/8
sensorDimension.path	https://cube.link/observedBy
timeDimension	http://odh.example.org/time-dimension/23/8
timeDimension.path	http://purl.org/dc/terms/date



Accessing the KG in BI tools

- The SQL connector emulates a PostgreSQL database
 - It works with its standard drivers (JDBC, ODBC, etc.)
- Works with popular tools
 - Tableau, PowerBI, Metabase (open-source)
 - Excel, Pandas, Veezoo

Let's demo it with Metabase (great at data exploration)



Q Search...

+ New



Browse data

Learn about our data

DATABASES > CUBE-ODH

classes

data_properties

object_properties

Performance and scalability

With TimescaleDB

- Demo made on a cheap 8 GB server with 4 cores (Hetzner CPX31)
- With 480M observations in the same hypertable
 - With default partitioning (7 days)
 - And indexing for quickly retrieving the values of a given time series
- Having the same data in a standard Postgres table on a large RDS instance was painful (both at ingestion and querying times)

Scaling with DuckDB

DuckDB

- “DuckDB is a fast in-process analytical database”
- “SQLite” for large tables/time series
- No daemon, JDBC to file

Process described in this post:

<https://www.linkedin.com/pulse/scaling-sparql-querying-billion-observations-ontop-duckdb-gschwend-myghf/>



Scaling with DuckDB: Setup

Synthetic CSV file

- 1 billion rows in observation table, ~41GB
- ~ 7'000 cities from Wikidata in city table (join)

Ontop 5.2.0 beta 2

- **No** keys in SQL schema for DuckDB (performance)
- R2RML file
- Ontop Lenses instead of keys
- MacBook M3 Pro (2023), 36 GB memory

Scaling with DuckDB: Results

Synthetic CSV file

- DuckDB load from CSV: A few minutes
- Written into an index file on disk

SPARQL queries

- Count
- AVG temp by city for a particular year

Ontop SPARQL endpoint

endpoint address: <http://localhost:8080/sparql> | ontop v5.2.0-SNAPSHOT

Observations x

S P O x

AVG x

Filter City x

Query x

Count x +

```
2 PREFIX cube: <https://cube.link/>
3 PREFIX rdf: <http://www.w3.org/1999/02/22-rdf-syntax-ns#>
4 PREFIX rdfs: <http://www.w3.org/2000/01/rdf-schema#>
5 PREFIX xsd: <http://www.w3.org/2001/XMLSchema#>
6 PREFIX dimension: <https://example.org/1b/dimension/>
7
8 SELECT ?cityName (AVG(?temp) as ?avgTemp) WHERE {
9     ?obs a cube:Observation ;
10         dimension:temperature ?temp ;
11         dimension:time ?time ;
12         dimension:city ?city .
13
14     ?city schema:name ?cityName .
15
16     FILTER(year(?time) = 2014)
17
18 }
19 GROUP BY ?cityName
20 ORDER BY ?cityName
21
```



Table

Response

Pivot Table

Google Chart

Geo



Ontop SPARQL endpoint

endpoint address: <http://localhost:8080/sparql> | ontop v5.2.0-SNAPSHOT

Observations x S P O x AVG x Filter City x Query x Query 1 x +

```
1 PREFIX schema: <http://schema.org/>
2 PREFIX cube: <https://cube.link/>
3 PREFIX rdf: <http://www.w3.org/1999/02/22-rdf-syntax-ns#>
4 PREFIX rdfs: <http://www.w3.org/2000/01/rdf-schema#>
5 PREFIX xsd: <http://www.w3.org/2001/XMLSchema#>
6 PREFIX dimension: <https://example.org/1b/dimension/>
7
8 SELECT ?cityName (AVG(?temp) as ?avgTemp) WHERE {
9     ?obs a cube:Observation ;
10         dimension:temperature ?temp ;
11         dimension:time ?time ;
12         dimension:city ?city .
13
14     ?city schema:name ?cityName .
```



Table

Response

Pivot Table

Google Chart

Geo

Showing 1 to 50 of 6,425 entries (in 7.054 seconds)

Search: Show 50 entries

	cityName	avgTemp
1	's-Hertogenbosch	"0.06506800886003021"^^xsd:decimal
2	6th of October City	"0.3289466888710063"^^xsd:decimal

Ontop SPARQL endpoint

endpoint address: <http://localhost:8080/sparql> | ontop v5.2.0-SNAPSHOT

Observations x S P O x AVG x Filter City x Query x Query 1 x +

```
3 PREFIX rdf: <http://www.w3.org/1999/02/22-rdf-syntax-ns#>
4 PREFIX rdfs: <http://www.w3.org/2000/01/rdf-schema#>
5 PREFIX xsd: <http://www.w3.org/2001/XMLSchema#>
6 PREFIX dimension: <https://example.org/1b/dimension/>
7
8 SELECT ?cityName (AVG(?temp) as ?avgTemp) WHERE {
9     ?obs a cube:Observation ;
10         dimension:temperature ?temp ;
11         dimension:time ?time ;
12         dimension:city ?city .
13
14     ?city schema:name ?cityName .
15
16     FILTER(year(?time) = 2014)
```



Table

Response

Pivot Table

Google Chart

Geo

Showing 1 to 50 of 6,425 entries (in 2.724 seconds)

Search: Show 50 entries

	cityName	avgTemp
1	's-Hertogenbosch	"0.3315944131998994"^^xsd:decimal
2	6th of October City	"0.038203605382636625"^^xsd:decimal

Ontop SPARQL endpoint

endpoint address: <http://localhost:8080/sparql> | ontop v5.2.0-SNAPSHOT

Observations x S P O x AVG x Filter City x Query x Count x +

```
1 PREFIX cube: <https://cube.link/>
2 PREFIX rdf: <http://www.w3.org/1999/02/22-rdf-syntax-ns#>
3 PREFIX rdfs: <http://www.w3.org/2000/01/rdf-schema#>
4 SELECT (COUNT(?sub) AS ?count) WHERE {
5     ?sub a cube:Observation .
6 }
7 LIMIT 10
```



Table

Response

Pivot Table

Google Chart

Geo

Showing 1 to 1 of 1 entries (in 0.191 seconds)

Search: Show 50 entries

count

1 "1000000000"^^xsd:integer

Showing 1 to 1 of 1 entries (in 0.191 seconds)

Scaling with DuckDB: Materialization

Same hardware, MacBook M3 Pro

- Ontop materialized the whole DB in 5h 15m
- 590GB N-Triples file
- Gzipped: 26GB
- 4'000'046'284 Triples

Triples per second: > **210'000!**



How about RDF triplestores?

- No known option that reaches this speed (DuckDB + Ontop)
- Discussions started with QLever team
- Check out presentation by Hannah Bast later today

Conclusions: Why RDF cubes?

Beyond CSV & co

- Clear semantics for **everyone else**
- Make implicit knowledge explicit
- Virtual SQL tables also benefits!

Data becomes discoverable!



Conclusions: Why RDF cubes?

RDF & SPARQL is on the web (or Intranet):

- Same dimensions can be truly shared
 - “Set of terms”
 - Taxonomies
- Relate concepts
 - owl:sameAs
 - Or any other relation that makes sense
- Virtual SQL tables also benefits!

Data becomes discoverable!



Thank you!

Questions?